



# Step 1 – Types of Databases



What all we'll learn today –

Simple – SQL vs NoSQL, how to create Postgres Databases, How to do CRUD on them

Advance – Relationships, Joins, Transactions

**There are a few types of databases, all serve different types of use-cases**

## NoSQL databases

1. Store data in a **schema-less** fashion. Extremely lean and fast way to store data.
2. Examples – MongoDB,



The screenshot shows a MongoDB interface with the following details:

- Collection:** SQL databases (3 of 11)
- Documents:** 20
- Total Size:** 2.0KB
- Avg. Size:** 104B
- Indexes:** 1
- Total Size:** 4.0KB
- Avg. Size:** 4.0KB

Document Preview:

```

_id: ObjectId("617604944e16573d5867039b")
name: "Seoul"
country: "South Korea"
continent: "Asia"
population: 25.674

_id: ObjectId("617604944e16573d5867039c")
name: "Mumbai"
country: "India"
continent: "Asia"
population: 19.98

_id: ObjectId("617604944e16573d5867039d")
name: "Lagos"
country: "Nigeria"
continent: "Africa"
population: 13.463

_id: ObjectId("617604944e16573d5867039e")

```

## Graph databases

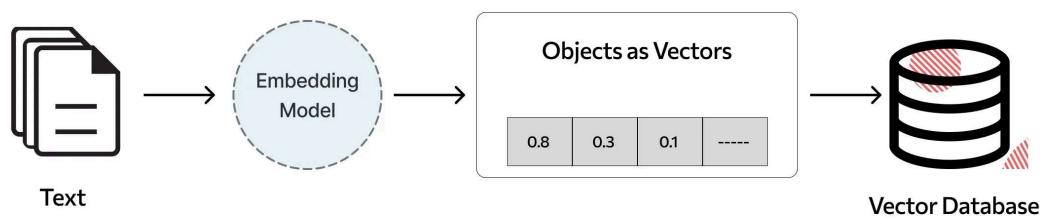
1. Data is stored in the form of a graph. Specially useful in cases where **relationships** need to be stored (social networks)
2. Examples – Neo4j





## Vector databases

1. Stores data in the form of vectors
2. Useful in Machine learning
3. Examples – Pinecone



graft

## SQL databases

1. Stores data in the form of rows
2. Most full stack applications will use this
3. Examples – MySQL, Postgres

# Step 2 – Why not NoSQL



SQL databases 3 of 11

You might've used MongoDB

It's **schemaless** properties make it ideal to for bootstrapping a project fast.

But as your app grows, this property makes it very easy for data to get **corrupted**

## What is schemaless?

Different rows can have different **schema** (keys/types)



```
1 _id: ObjectId('65b3f3319e01786d275501c8')
2 userId: 65b3f3319e01786d275501c6
3 balance: 5885.875967139374
4 __v: 0
```



```
_id: ObjectId('65b3f3469e01786d275501ce')
userId: ObjectId('65b3f3469e01786d275501cc')
__v: 0
amountBalance: 720.3759487457523
```



```
_id: ObjectId('65b3f3499e01786d275501d3')
userId: ObjectId('65b3f3499e01786d275501d1')
balance: "harkirat"
__v: 0
```

## Problems?

1. Can lead to inconsistent database
2. Can cause runtime errors
3. Is too flexible for an app that needs strictness

Up Next:

1. Can move very fast



2 SQL databases 3 of 11

ma very easily



You might think that `mongoose` does add strictness to the codebase because we used to define a schema there. That strictness is present at the Node.js level, not at the DB level. You can still put in erroneous data in the database that doesn't follow that schema.

## Step 3 – Why SQL?

SQL databases have a strict schema. They require you to

1. Define your schema
2. Put in data that follows that schema
3. Update the schema as your app changes and perform migrations

So there are 4 parts when using an SQL database (not connecting it to Node.js, just running it and putting data in it)

1. Running the database.
2. Using a library that lets you connect and put data in it.

's schema .

4. Run queries on the database to interact with the data



SQL databases 3 of 11 (Delete)

## Step 4 - Creating a database

You can start a Postgres database in a few ways -

### ▼ Using neondb

<https://neon.tech/> is a decent service that lets you create a server.

SQL databases 3 of 11

Name	Region	Created at	Storage	Postgres version	⋮
test	US East (Ohio)	Sep 3, 2023 4:23 pm	99 MiB	15	⋮

**Connection Details**

Branch: main (PRIMARY) Compute: RW (ACTIVE)

Database: calm-gobbler-41.db.2253874 Role: 100xdevs Reset password

Connection string: postgresql://100xdevs:\*\*\*\*\*@ep-broken-frost-69135494.us-east-2.ows.neon.tech/calm-gobbler-41.db.2253874?sslmode=require

Your password is saved in a secure storage vault. More about [connecting from different languages, frameworks, and platforms](#).

**Operations**

**Branches**

Name: main (PRIMARY) RW compute Used space: 0.25 CU (ACTIVE) 1%

**Usage since January 1**

Active time all computes: 0% 0 h / 100 h

The Free Tier grants 100 hours of Active time per month for all computes. Exceeding this limit, non-primary branch computes are subject to suspension. The primary branch compute remains available.

To increase project limits, [upgrade to Pro](#)

**Storage**

Project storage: 99 MiB

Branches: 1 / 10

History retention: 7 days

Current data size for main (PRIMARY): 49 MiB / 3 GiB

## Connection String

postgresql://username:password@ep-broken-frost-69135494

### ▼ Using docker locally

docker run --name my-postgres -e POSTGRES\_PASSWORD=m

## Connection String

postgresql://postgres:mysecretpassword@localhost:5432/pos

## How to run PostgreSQL in windows terminal(if you have docker)



### SQL databases 3 of 11

- First run a Docker GUI application that helps in running commands in terminal.
- After that run it with the Docker instance by the help of following command .
  - for the first time if the image is not downloaded .
  - `docker run --name my-postgres1 -e POSTGRES_PASSWORD=mysecretpassword -d -p 5432:5432 postgres` .
  - if the Docker image is there, prior to use it can simply be runned by `docker run <image name>` .
- After that ,
  - use `docker exec -it my-postgres1 psql -U postgres -d postgres` this command in terminal .
  - then enter the password and it will connect to localhost Postgres instance .
  - now you will be inside the PostgreSQL command line that looks like `postgres#` .
  - You can check it by running `\dt` , (the command to display all the tables.)



The connection string is similar to the string we had in Mongoose.

## Connection String

`postgresql://username:password@host/database`

Where does Harkirat live  $\Rightarrow [1, 2, 2, 2, 3001, 100]$

Harkirat lives in India  $\Rightarrow [1, 2, 2, 2, 2, 2]$

≡ SQL databases 3 of 11

# Using a library that let's you connect and put data in it.

## 1. psql

`psql` is a terminal-based front-end to PostgreSQL. It provides an interactive command-line interface to the PostgreSQL (or TimescaleDB) database. With `psql`, you can type in queries interactively, issue them to PostgreSQL, and see the query results.

### How to connect to your database?

`psql` Comes bundled with `postgresql`. You don't need it for this tutorial. We will directly be communicating with the database from `Node.js`.

```
psql -h p-broken-frost-69135494.us-east-2.aws.neon.tech -d dat
```



## 2. pg

`pg` is a `Node.js` library that you can use in your backend app to store data in the Postgres DB (similar to `mongoose`). We will be installing this eventually in our app.

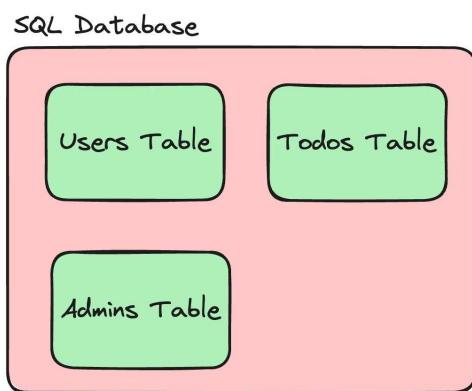


SQL databases 3 of 11

# Creating a table and defining its schema .

## Tables in SQL

A single database can have multiple tables inside. Think of them as collections in a MongoDB database.



Until now, we have a database that we can interact with. The next step in case of postgres is to define the **schema** of your tables.

SQL stands for **Structured query language**. It is a language in which you can describe what/how you want to put data in the database.

To create a table, the command to run is -

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    username VARCHAR(50) UNIQUE NOT NULL,
    email VARCHAR(255) UNIQUE NOT NULL,
    password VARCHAR(255) NOT NULL,
```



```
created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
```



SQL databases 3 of 11

There are a few parts of this SQL statement, let's decode them one by one

## 1. CREATE TABLE users

**CREATE TABLE users** : This command initiates the creation of a new table in the database named **users**.

## 2. id SERIAL PRIMARY KEY

- **id** : The name of the first column in the **users** table, typically used as a unique identifier for each row (user). Similar to **\_id** in mongoDB
- **SERIAL** : A PostgreSQL-specific data type for creating an auto-incrementing integer. Every time a new row is inserted, this value automatically increments, ensuring each user has a unique **id**.
- **PRIMARY KEY** : This constraint specifies that the **id** column is the primary key for the table, meaning it uniquely identifies each row. Values in this column must be unique and not null.

## 3. email VARCHAR(255) UNIQUE NOT NULL,

- **email** : The name of the second column, intended to store the user's username.
- **VARCHAR(50)** : A variable character string data type that can store up to 50 characters. It's used here to limit the length of the username.
- **UNIQUE** : This constraint ensures that all values in the **username** column are unique across the table. No two users can have the same username.

nts null values from being inserted into the **username** column. Every row must have a

username value.



SQL databases 3 of 11

## 4. password VARCHAR(255) NOT NULL

Same as above, can be non unique

## 5. created\_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT\_TIMESTAMP

- **created\_at** : The name of the fifth column, intended to store the timestamp when the user was created.
- **TIMESTAMP WITH TIME ZONE** : This data type stores both a timestamp and a time zone, allowing for the precise tracking of when an event occurred, regardless of the user's or server's time zone.
- **DEFAULT CURRENT\_TIMESTAMP** : This default value automatically sets the **created\_at** column to the date and time at which the row is inserted into the table, using the current timestamp of the database server.



If you have access to a database right now, try running this command to create a simple table in there

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    username VARCHAR(50) UNIQUE NOT NULL,
    email VARCHAR(255) UNIQUE NOT NULL,
    password VARCHAR(255) NOT NULL,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT
    CURRENT_TIMESTAMP
);
```

Then try running

\dt+

reated or not



# Step 7 – Interacting with the database

There are 4 things you'd like to do with a database

## 1. INSERT

```
INSERT INTO users (username, email, password)   
VALUES ('username_here', 'user@example.com', 'user_password');
```

 Notice how you didn't have to specify the `id` because it auto increments

## 2. UPDATE

```
UPDATE users   
SET password = 'new_password'  
WHERE email = 'user@example.com';
```

### 3. DELETE



SQL databases 3 of 11



```
WHERE id = 1;
```

### 4. Select

```
SELECT * FROM users  
WHERE id = 1;
```



Try running all 4 of these in your terminal if you have `psql` installed locally.

If not, that's fine we'll eventually be doing these through the `pg` library

## Step 8 – How to do queries from a Node.js app?

In the end, postgres exposes a protocol that someone needs to talk to be able to send these commands (update, delete) to the database.

`psql` is one such library that takes commands from your database.

To do the same in a Node.js , you can use one of many [Postgres](#)



[SQL databases](#) 3 of 11

## pg library

<https://www.npmjs.com/package/pg>

Non-blocking PostgreSQL client for Node.js.

Documentation - <https://node-postgres.com/>

Connecting -

```
import { Client } from 'pg'

const client = new Client({
  host: 'my.database-server.com',
  port: 5334,
  database: 'database-name',
  user: 'database-user',
  password: 'secretpassword!!'
})

client.connect()
```



Querying -

```
const result = await client.query('SELECT * FROM USERS;')
console.log(result)
```



// write a function to create a users table in your database.

```
import { Client } from 'pg'
```

```
const client = new Client({
  connectionString: "postgresql://postgres:mysecretpassword@loc
})
```



```
const result = await client.query(`  
  users (  
    MARY KEY,  
    username VARCHAR(50) UNIQUE NOT NULL,  
    email VARCHAR(255) UNIQUE NOT NULL,  
    password VARCHAR(255) NOT NULL,  
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_T  
  );  
)  
  console.log(result)  
}  
  
createUsersTable();
```

# Step 9 - Creating a simple Node.js app

1. Initialise an empty typescript project

```
npm init -y  
npx tsc --init
```



1. Change the `rootDir` and `outDir` in `tsconfig.json`



1. Install the `pg` library and its types (because we're using TS)

三

SQL databases 3 of 11



npm install @types/pg



1. Create a simple Node.js app that lets you put data

Create a function that lets you insert data into a table. Make it `async`, make sure `client.connect` resolves before you do the insert.

▼ Answer

```
import { Client } from 'pg';
```

2

```
// Async function to insert data into a table
async function insertData() {
  const client = new Client({
    host: 'localhost',
    port: 5432,
    database: 'postgres',
    user: 'postgres',
    password: 'mysecretpassword',
  });

  try {
    await client.connect(); // Ensure client connection is established
    const insertQuery = "INSERT INTO users (username, email, password) VALUES ('john_doe', 'john@example.com', 'hashed_password')";
    const res = await client.query(insertQuery);
    console.log('Insertion success:', res); // Output insertion result
  } catch (err) {
    console.error('Error during the insertion:', err);
  } finally {
    await client.end(); // Close the client connection
  }
}

insertData();
```



This is an **insecure** way to store data in your tables.

...nality eventually via [HTTP](#),

someone can do an **SQL INJECTION** to get access to your **SQL databases** 3 of 11 ur data.

1. More secure way to store data.

Update the code so you don't put **user provided fields** in the **SQL string**

▼ What are user provided strings?

In your final app, this insert statement will be done when a user signs up on your app.

Email, username, password are all user provided strings

▼ What is the **SQL string** ?

```
const insertQuery = "INSERT INTO users (username, email, pass
```

◀ ▶

▼ Hint

```
const insertQuery = 'INSERT INTO example_table(column1, colu
```

◀ ▶

▼ Solution

```
import { Client } from 'pg';
```

copy

```
// Async function to insert data into a table
async function insertData(username: string, email: string, passw
  const client = new Client({
    host: 'localhost',
    port: 5432,
    database: 'postgres',
    user: 'postgres',
    password: 'mysecretpassword',
  });

  try {
    await client.connect(); // Ensure client connection is establis
```

to prevent SQL injection  
ITO users (username, email, pas  
const values = [username, email, password];

```

    const res = await client.query(insertQuery, values);
    insertionSuccess(res); // Output insertion result

    console.error('Error during the insertion:', err);
} finally {
    await client.end(); // Close the client connection
}
}

// Example usage
insertData('username5', 'user5@example.com', 'user_password');

```

## 1. Query data

Write a function `getUser` that lets you fetch data from the database given a `email` as input.

```

import { Client } from 'pg';

// Async function to fetch user data from the database given an email
async function getUser(email: string) {
    const client = new Client({
        host: 'localhost',
        port: 5432,
        database: 'postgres',
        user: 'postgres',
        password: 'mysecretpassword',
    });

    try {
        await client.connect(); // Ensure client connection is established
        const query = 'SELECT * FROM users WHERE email = $1';
        const values = [email];
        const result = await client.query(query, values);

        if (result.rows.length > 0) {
            console.log('User found:', result.rows[0]); // Output user data
            return result.rows[0]; // Return the user data
        } else {
            console.log(`No user found for the given email.`);
            return null;
        }
    } catch (err) {
        console.error('Error fetching user data:', err);
    } finally {
        client.end();
    }
}

```

```
        } catch (err) {  
            // Throw or handle error appropriately  
            throw err;  
        } finally {  
            await client.end(); // Close the client connection  
        }  
  
    }  
  
    // Example usage  
    getUser('user5@example.com').catch(console.error);
```

# Step 10 – Relationships and Transactions

Relationships let you store data in different tables and relate it with each other.

## Relationships in Mongodb

Since `mongodb` is a NoSQL database, you can store any shape of data in it.

If I ask you to store a user's details along with their address, you can store it in an object that has the address details.

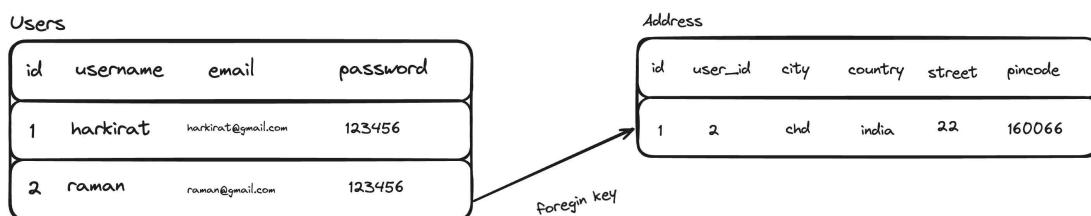
SQL databases 3 of 11 [tId\('65b3f3469e01786d275501ce'\)](#)

[Address Object](#)

```
street: "1 West HW"
city: "Chandigarh"
country: "India"
pincode: "160066"
email: "harkirat@gmail.com"
name: "harkirat"
```

## Relationships in SQL

Since SQL can not store **objects** as such, we need to define two different tables to store this data in.



This is called a **relationship**, which means that the **Address** table is related to the **Users** table.

When defining the table, you need to define the **relationship**

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    username VARCHAR(50) UNIQUE NOT NULL,
    email VARCHAR(255) UNIQUE NOT NULL,
    password VARCHAR(255) NOT NULL,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
```

```
CREATE TABLE addresses (
    id PRIMARY KEY,
    user_id NOT NULL,
    city VARCHAR(100) NOT NULL,
    country VARCHAR(100) NOT NULL,
    street VARCHAR(255) NOT NULL,
    pincode VARCHAR(20),
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP,
    FOREIGN KEY (user_id) REFERENCES users(id) ON DELETE CASCADE
);
```

## SQL query

To insert the address of a user –

```
INSERT INTO addresses (user_id, city, country, street, pincode)
VALUES (1, 'New York', 'USA', '123 Broadway St', '10001');
```

Now if you want to get the address of a user given an `id` , you can run the following query –

```
SELECT city, country, street, pincode
FROM addresses
WHERE user_id = 1;
```

## Extra – Transactions in SQL

 Good question to have at this point is what queries are run when the user signs up and sends both their information and their address in a single request. Do we send two SQL queries into the database? What if one of the queries (address query for example) fails? This would require `transactions` in SQL to ensure either

the address goes in, or neither

## ▼ SQL Query

SQL databases 3 of 11

ransaction

```
INSERT INTO users (username, email, password)
VALUES ('john_doe', 'john_doe@example.com', 'securepassword')
```

```
INSERT INTO addresses (user_id, city, country, street, pincode)
VALUES (currval('users_id_seq'), 'New York', 'USA', '123 Broadway')
```

## COMMIT;

◀ ▶

## ▼ Node.js Code

```
import { Client } from 'pg';
```

2

```
async function insertUserAndAddress(  
    username: string,  
    email: string,  
    password: string,  
    city: string,  
    country: string,  
    street: string,  
    pincode: string  
) {  
    const client = new Client({  
        host: 'localhost',  
        port: 5432,  
        database: 'postgres',  
        user: 'postgres',  
        password: 'mysecretpassword',  
    });
```

```
try {
```

```
await client.connect();
```

```
// Start transaction
```

```
await client.query('BEGIN');
```

// Insert user

ame, email, password)

## VALUES (\$1, \$2, \$3)

RETURNING id;



SQL databases 3 of 11

```
res = await client.query(insertUserText, [username, password, city, country, street, pincode]);
const userId = userRes.rows[0].id;

// Insert address using the returned user ID
const insertAddressText = `
  INSERT INTO addresses (user_id, city, country, street, pincode)
  VALUES ($1, $2, $3, $4, $5);
`;
await client.query(insertAddressText, [userId, city, country, street, pincode]);

// Commit transaction
await client.query('COMMIT');

console.log('User and address inserted successfully');
} catch (err) {
  await client.query('ROLLBACK'); // Roll back the transaction
  console.error('Error during transaction, rolled back.', err);
  throw err;
} finally {
  await client.end(); // Close the client connection
}
}

// Example usage
insertUserAndAddress(
  'johndoe',
  'john.doe@example.com',
  'securepassword123',
  'New York',
  'USA',
  '123 Broadway St',
  '10001'
);
```

# Step 11 – Joins



SQL databases 3 of 11

Defining relationships is easy.

What's hard is **joining** data from two (or more) tables together.

For example, if I ask you to fetch me a user's details **and** their address, what SQL would you run?

## ▼ Approach 1 (Bad)

-- Query 1: Fetch user's details



```
SELECT id, username, email
FROM users
WHERE id = YOUR_USER_ID;
```

-- Query 2: Fetch user's address



```
SELECT city, country, street, pincode
FROM addresses
WHERE user_id = YOUR_USER_ID;
```

## ▼ Approach 2 (using joins)

```
SELECT users.id, users.username, users.email, addresses.city, a.c
FROM users
JOIN addresses ON users.id = addresses.user_id
WHERE users.id = '1';
```



```
SELECT u.id, u.username, u.email, a.city, a.country, a.street, a.p
FROM users u
JOIN addresses a ON u.id = a.user_id
WHERE u.id = YOUR_USER_ID;
```



Now try converting the same to your node app

## ▼ Approach 1 (Bad)

```
import { Client } from 'pg';
```



details and address separately

```
async function getUserDetailsAndAddressSeparately(userId: sti
```



SQL databases 3 of 11

```

const client = new Client({
  host: 'localhost',
  database: 'postgres',
  user: 'postgres',
  password: 'mysecretpassword',
});

try {
  await client.connect();

  // Fetch user details
  const userDetailsQuery = 'SELECT id, username, email FROM users';
  const userDetails = await client.query(userDetailsQuery, [userId]);

  // Fetch user address
  const userAddressQuery = 'SELECT city, country, street, pincode FROM addresses WHERE user_id = $1';
  const userAddress = await client.query(userAddressQuery, [userId]);

  if (userDetails.rows.length > 0) {
    console.log('User found:', userDetails.rows[0]);
    console.log('Address:', userAddress.rows.length > 0 ? userAddress.rows[0] : null);
    return { user: userDetails.rows[0], address: userAddress.rows[0] };
  } else {
    console.log('No user found with the given ID.');
    return null;
  }
} catch (err) {
  console.error('Error during fetching user and address:', err);
  throw err;
} finally {
  await client.end();
}
}

getUserDetailsAndAddressSeparately("1");

```

## ▼ Approach 2 (using joins)

```
import { Client } from 'pg';
```



```

data and their address together
withAddress(userId: string) {
  const client = new Client({

```

```
host: 'localhost',
  SQL databases 3 of 11
    postgres',
    user: 'postgres',
    password: 'mysecretpassword',
  );
```

```
try {
  await client.connect();
  const query = `
    SELECT u.id, u.username, u.email, a.city, a.country, a.street
    FROM users u
    JOIN addresses a ON u.id = a.user_id
    WHERE u.id = $1
  `;
  const result = await client.query(query, [userId]);

  if (result.rows.length > 0) {
    console.log('User and address found:', result.rows[0]);
    return result.rows[0];
  } else {
    console.log('No user or address found with the given ID.')
    return null;
  }
} catch (err) {
  console.error('Error during fetching user and address:', err);
  throw err;
} finally {
  await client.end();
}
}

getUserDetailsWithAddress("1");
```

No

## Benefits of using a join -

## 2. Simplified Application Logic



3 SQL databases 3 of 11

egrity

# Types of Joins

## 1. INNER JOIN

Returns rows when there is at least one match in both tables. If there is no match, the rows are not returned. It's the most common type of join.

**Use Case: Find All Users With Their Addresses. If a user hasn't filled their address, that user shouldn't be returned**

```
SELECT users.username, addresses.city, addresses.country, address
FROM users
INNER JOIN addresses ON users.id = addresses.user_id;
```



## 2. LEFT JOIN

Returns all rows from the left table, and the matched rows from the right table.

Use case - To list all users from your database along with their address information (if they've provided it), you'd use a LEFT JOIN. Users without an address will still appear in your query result, but the address fields will be NULL for them.

```
SELECT users.username, addresses.city, addresses.country, address
FROM users
LEFT JOIN addresses ON users.id = addresses.user_id;
```



## 3. RIGHT JOIN

Returns all rows from the right table, and the matched rows from the left table.

... (continued)

Use case - Given the structure of the database, a RIGHT JOIN

- ☰ **SQL databases** 3 of 11 on since the **addresses** table is unlikely to have entries not linked to a user due to the foreign key constraint. However, if you had a situation where you start with the **addresses** table and optionally include user information, this would be the theoretical use case.

```
SELECT users.username, addresses.city, addresses.country, address
FROM users
RIGHT JOIN addresses ON users.id = addresses.user_id;
```



## 4. FULL JOIN

Returns rows when there is a match in one of the tables. It effectively combines the results of both LEFT JOIN and RIGHT JOIN.

Use case - A FULL JOIN would combine all records from both **users** and **addresses**, showing the relationship where it exists. Given the constraints, this might not be as relevant because every address should be linked to a user, but if there were somehow orphaned records on either side, this query would reveal them.

```
SELECT users.username, addresses.city, addresses.country, address
FROM users
FULL JOIN addresses ON users.id = addresses.user_id;
```

